

· How We Scaled Our Python Backend to 10M Requests / Day

Most teams optimize queries way too early.

Here is what actually made the difference for our infrastructure:

1. Redis In-Memory Caching on Hot Paths
2. Async Connection Pooling for PostgreSQL
3. Background Worker Decoupling via Celery

The result? Latency dropped from 380ms down to 14ms.

Drop a comment below if you want the open-source benchmark repo!